

ALGORITHM FOR POD SCHEDULING IN KUBERNETES

INPUT:

- **PodsQueue**: Queue of unscheduled pods
- **Nodes**: Set of available nodes
- **PodsTypeInfo**: Information about pod types and their affinities/anti-affinities
- **ResourceInfo**: Resources available on nodes (CPU, memory, etc.)
- **TrafficMatrix**: Matrix representing inter-pod communication traffic
- **LatencyMatrix**: Matrix representing network latency between nodes
- **CostConversionFactors**: Factors to convert resource usage into cost-equivalent terms (e.g., CPU cost/unit, memory cost/unit)

STEPS

1. Initialization

- Compute **PodsTraffic** from **TrafficMatrix**, where **PodsTraffic[i][j]** represents the communication traffic between pod *i* and pod *j*.
- Extract resource requirements and affinity/anti-affinity rules for each pod from **PodsQueue**.

2. Preprocess Nodes and Pods

- Group nodes by resource capacity (e.g., high-memory, high-CPU, balanced).
- Categorize pods into workload types (e.g., compute-intensive, memory-intensive, balanced).

3. Convert Resource Usage into Cost

For each node *N*:

- Compute the **cost equivalence** of available resources:

$$\text{CostAvailable}(N) = \text{CPU}(N) \cdot \text{CPUFactor} + \text{Memory}(N) \cdot \text{MemoryFactor}$$

- Similarly, for each pod *P*, compute the **resource request cost**:

$$\text{CostRequested}(P) = \text{CPU}(P) \cdot \text{CPUFactor} + \text{Memory}(P) \cdot \text{MemoryFactor}$$

4. Calculate Node Scores for Each Pod

For each pod P in **PodsQueue**:

- For each node N in **Nodes**:

1. **Compute Resource Cost Score**

$$\text{ResourceCostScore}(P, N) = \text{CostAvailable}(N) - \text{CostRequested}(P)$$

2. **Compute Latency and Traffic Cost**

- For already scheduled pods Q on N :

$$\text{TrafficCost}(P, Q) = \text{PodsTraffic}[P][Q] \cdot \text{LatencyMatrix}[N][\text{Node}(Q)]$$

- Sum the traffic cost:

$$\text{TrafficCostScore}(P, N) = \sum_{Q \in \text{PodsOn}(N)} \text{TrafficCost}(P, Q)$$

3. **Compute Final Node Score**

$$\text{NodeScore}(P, N) = \alpha \cdot \text{ResourceCostScore} - \beta \cdot \text{TrafficCostScore}$$

(where α, β are weight coefficients).

5. Assign Pods to Nodes

- While **PodsQueue** is not empty:
 - Pop the highest-priority pod P from **PodsQueue**.
 - Sort nodes N by their **NodeScore** for P in descending order.
 - Assign P to the highest-scoring node N that satisfies all constraints:
 - Sufficient resources.
 - No anti-affinity violations.
 - Affinity rules are satisfied.
 - Update resource availability and node state.
 - Recompute **PodsTraffic** and **LatencyMatrix** scores for remaining pods if necessary.

7. Return Result

- Output the final pod-to-node assignments.
-

MODIFICATIONS SUGGESTED

1. Consider pod qos class type to add separate score for pods.
2. Consider deployments with hpa and vpa .
3. In step 3, consider the actual pod level usage.